

Computer Vision HW2 Report

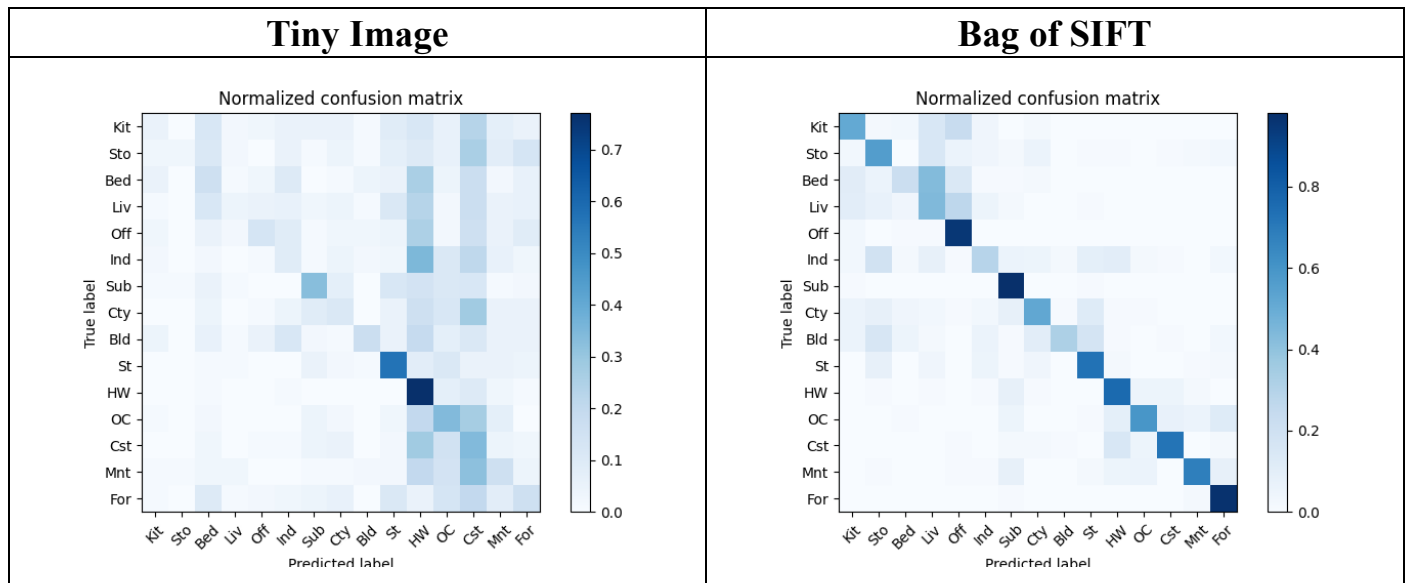
Student ID: R12K41025

Name: 杜冠廷

Part 1. (10%)

- Plot confusion matrix of two settings. (i.e. Bag of sift and tiny image) (5%)

Ans:



- Compare the results/accuracy of both settings and explain the result. (5%)

Ans:

Accuracy	
Tiny Image	0.23
Bag of SIFT	0.616

As expected, the performance of Tiny Image is worse than that of Bag of SIFT because the Tiny Image approach can only capture very basic image features, such as color distribution and brightness. In contrast, Bag of SIFT can extract more detailed and discriminative features, such as textures, edges, and local shapes. The process of building the vocabulary and generating the vocab.pkl file takes a significant amount of time. To reduce this, we limited the vocabulary construction to only 1/10 of the images from the training set and extracted a maximum of 80 feature points per image.

Part 2. (25%)

- Report accuracy of both models on the validation set. (2%)

Ans:

Accuracy	
MyNet	0.7464
ResNet18	0.9234

- Print the network architecture & number of parameters of both models. What is the main difference between ResNet and other CNN architectures? (5%)

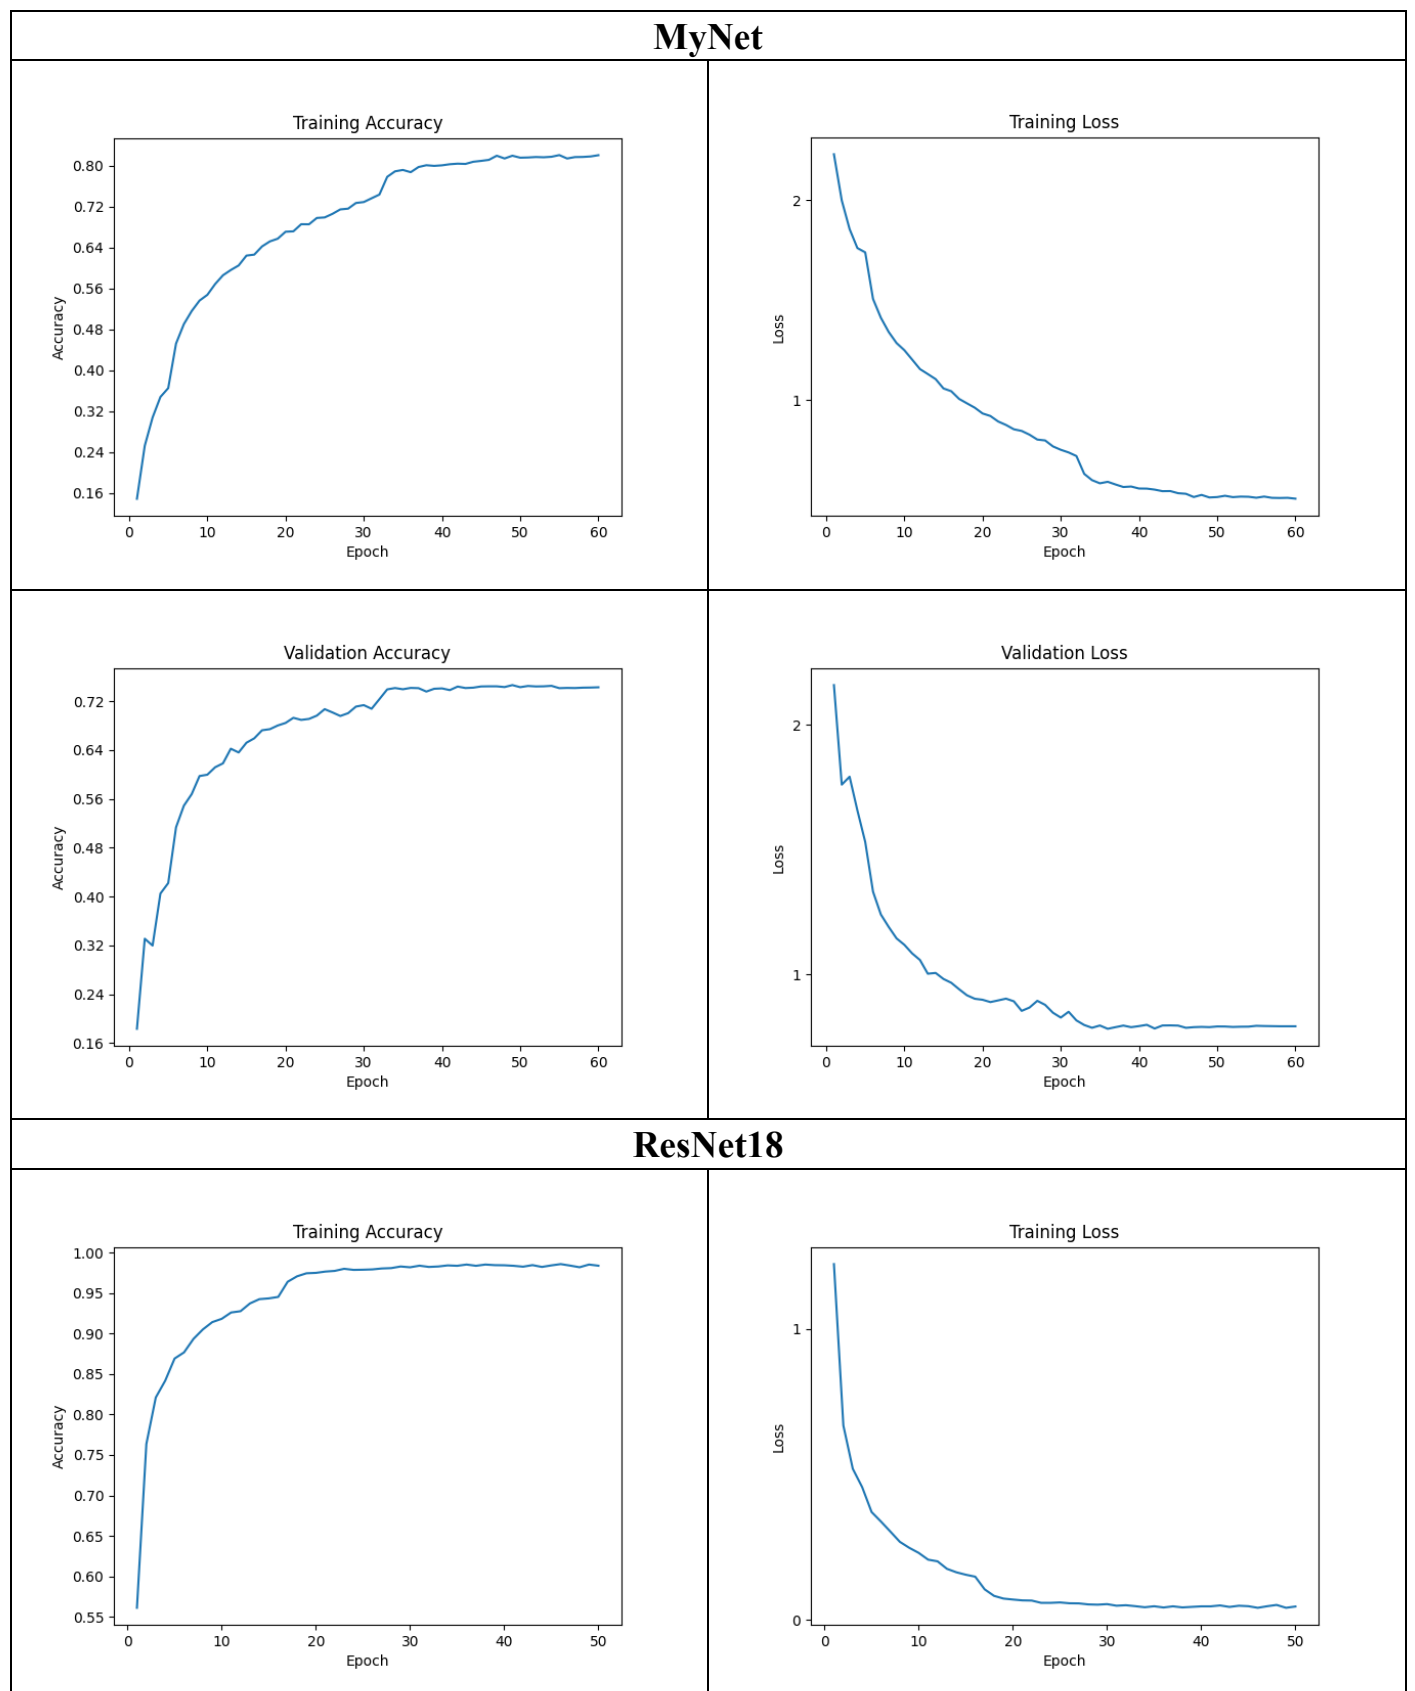
Ans:


MyNet		
Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 224, 224]	1,792
ReLU-2	[-1, 64, 224, 224]	0
Conv2d-3	[-1, 64, 224, 224]	36,928
ReLU-4	[-1, 64, 224, 224]	0
MaxPool2d-5	[-1, 64, 112, 112]	0
Conv2d-6	[-1, 128, 112, 112]	73,856
ReLU-7	[-1, 128, 112, 112]	0
Conv2d-8	[-1, 128, 112, 112]	147,584
ReLU-9	[-1, 128, 112, 112]	0
MaxPool2d-10	[-1, 128, 28, 28]	0
Conv2d-11	[-1, 256, 28, 28]	295,168
ReLU-12	[-1, 256, 28, 28]	0
Conv2d-13	[-1, 256, 28, 28]	590,080
ReLU-14	[-1, 256, 28, 28]	0
MaxPool2d-15	[-1, 256, 7, 7]	0
Flatten-16	[-1, 12544]	0
Linear-17	[-1, 1024]	12,846,080
ReLU-18	[-1, 1024]	0
Linear-19	[-1, 512]	524,800
ReLU-20	[-1, 512]	0
Linear-21	[-1, 10]	5,130
Total params: 14,521,418		
Trainable params: 14,521,418		
Non-trainable params: 0		
Input size (MB): 0.57		
Forward/backward pass size (MB): 160.23		
Params size (MB): 55.39		
Estimated Total Size (MB): 216.20		

ResNet18		
Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 114, 114]	1,728
BatchNorm2d-2	[-1, 64, 114, 114]	128
ReLU-3	[-1, 64, 114, 114]	0
MaxPool2d-4	[-1, 64, 57, 57]	0
Conv2d-5	[-1, 64, 57, 57]	36,864
BatchNorm2d-6	[-1, 64, 57, 57]	128
ReLU-7	[-1, 64, 57, 57]	0
Conv2d-8	[-1, 64, 57, 57]	36,864
BatchNorm2d-9	[-1, 64, 57, 57]	128
ReLU-10	[-1, 64, 57, 57]	0
BasicBlock-11	[-1, 64, 57, 57]	0
Conv2d-12	[-1, 64, 57, 57]	36,864
BatchNorm2d-13	[-1, 64, 57, 57]	128
ReLU-14	[-1, 64, 57, 57]	0
Conv2d-15	[-1, 64, 57, 57]	36,864
BatchNorm2d-16	[-1, 64, 57, 57]	128
ReLU-17	[-1, 64, 57, 57]	0
BasicBlock-18	[-1, 64, 57, 57]	0
Conv2d-19	[-1, 128, 29, 29]	73,728
BatchNorm2d-20	[-1, 128, 29, 29]	256
ReLU-21	[-1, 128, 29, 29]	0
Conv2d-22	[-1, 128, 29, 29]	147,456
BatchNorm2d-23	[-1, 128, 29, 29]	256
Conv2d-24	[-1, 128, 29, 29]	8,192
BatchNorm2d-25	[-1, 128, 29, 29]	256
ReLU-26	[-1, 128, 29, 29]	0
BasicBlock-27	[-1, 128, 29, 29]	0
Conv2d-28	[-1, 128, 29, 29]	147,456
BatchNorm2d-29	[-1, 128, 29, 29]	256
ReLU-30	[-1, 128, 29, 29]	0
Conv2d-31	[-1, 128, 29, 29]	147,456
BatchNorm2d-32	[-1, 128, 29, 29]	256
ReLU-33	[-1, 128, 29, 29]	0
BasicBlock-34	[-1, 128, 29, 29]	0
Conv2d-35	[-1, 256, 15, 15]	294,912
BatchNorm2d-36	[-1, 256, 15, 15]	512
ReLU-37	[-1, 256, 15, 15]	0
Conv2d-38	[-1, 256, 15, 15]	589,824
BatchNorm2d-39	[-1, 256, 15, 15]	512
Conv2d-40	[-1, 256, 15, 15]	32,768
BatchNorm2d-41	[-1, 256, 15, 15]	512
ReLU-42	[-1, 256, 15, 15]	0
BasicBlock-43	[-1, 256, 15, 15]	0
Conv2d-44	[-1, 256, 15, 15]	589,824
BatchNorm2d-45	[-1, 256, 15, 15]	512
ReLU-46	[-1, 256, 15, 15]	0
Conv2d-47	[-1, 256, 15, 15]	589,824
BatchNorm2d-48	[-1, 256, 15, 15]	512
ReLU-49	[-1, 256, 15, 15]	0
BasicBlock-50	[-1, 256, 15, 15]	0
Conv2d-51	[-1, 512, 8, 8]	1,179,648
BatchNorm2d-52	[-1, 512, 8, 8]	1,024
ReLU-53	[-1, 512, 8, 8]	0
Conv2d-54	[-1, 512, 8, 8]	2,359,296
BatchNorm2d-55	[-1, 512, 8, 8]	1,024
Conv2d-56	[-1, 512, 8, 8]	131,072
BatchNorm2d-57	[-1, 512, 8, 8]	1,024
ReLU-58	[-1, 512, 8, 8]	0
BasicBlock-59	[-1, 512, 8, 8]	0
Conv2d-60	[-1, 512, 8, 8]	2,359,296
BatchNorm2d-61	[-1, 512, 8, 8]	1,024
ReLU-62	[-1, 512, 8, 8]	0
Conv2d-63	[-1, 512, 8, 8]	2,359,296
BatchNorm2d-64	[-1, 512, 8, 8]	1,024
ReLU-65	[-1, 512, 8, 8]	0
BasicBlock-66	[-1, 512, 8, 8]	0
AdaptiveAvgPool2d-67	[-1, 512, 1, 1]	0
Linear-68	[-1, 10]	5,130
ResNet-69	[-1, 10]	0
Total params: 11,173,962		
Trainable params: 11,173,962		
Non-trainable params: 0		
Input size (MB): 0.57		
Forward/backward pass size (MB): 67.01		
Params size (MB): 42.63		
Estimated Total Size (MB): 110.21		

• Plot four learning curves (loss & accuracy) of the training process (train/validation) for both models. Total 8 plots. (8%)

Ans:





• Briefly describe what method do you apply on your best model? (e.g. data augmentation, model architecture, loss function, etc) (10%)

Ans:

Since ResNet18 is designed for classifying the large-scale dataset, I resized the images on CIFAR-10 to 224×224 to prevent excessive information loss caused by the Max Pooling layer. Additionally, due to the limited amount of training data, I applied several data augmentation techniques to prevent overfitting:

1. `transforms.RandomResizedCrop(224, scale=(256/480, 1.0))`
2. `transforms.RandomHorizontalFlip()`
3. `transforms.RandomRotation(10)`
4. `transforms.RandomGrayscale()`